# Helion vs selected baselines on GB300

Forward benchmarks · 14 kernels · geometric mean across six shapes · 168/168 recorded comparisons passed

† CuTe requests include Triton stages or fallbacks for K6. K3 now uses CuTe throughout. Details are in the JSON receipts.

Original _steady_bench · CUDA graphs · five runs · cold L2 · K16 β=0.1 · 1400 W GPU limit · Helion bfc25902 + CuTe patches (K3, K4, K6, K19)

K3 updated on GPU 3 (case 1: M8; 11-file CuTe patch 18a14109); other 13 kernels retain historical results.